

# Computer Architecture

Lecture 04

Interconnection  
Structure

# Reference Books

- Computer Organization and Architecture:  
Designing for Performance-William  
Stallings  
(8<sup>th</sup> Edition)
  - Any later edition is fine

# Overview

- The collection of paths connecting the various modules is called the *interconnection structure*
- The interconnection structure must support the following types of transfers:
  - **Memory to processor:** The processor reads an instruction or a unit of data from memory
  - **Processor to memory:** The processor writes a unit of data to memory
  - **I/O to processor:** The processor reads data from an I/O device via an I/O module
  - **Processor to I/O:** The processor sends data to the I/O device
  - **I/O to or from memory:** For these two cases, an I/O module is allowed to exchange data directly with memory, without involving the processor, in direct memory access (DMA)

# BUS Interconnection

- A bus is a communication pathway connecting two or more devices
  - It is a **shared transmission medium**
  - Consists of **multiple** communication pathways, or **lines**, each line is capable of transmitting signals representing binary 1 and binary 0
- Multiple devices connect to the bus, and a signal transmitted by any one device is available for reception by all other devices attached to the bus
- If two devices transmit during the same time period, their signals will overlap and become garbled
  - Only one device at a time can successfully transmit
- A bus that connects major computer components (processor, memory, I/O) is called a **system bus**

# Bus Structure

- A system bus consists, typically, of from about 50 to hundreds of separate lines
- Each line is assigned a particular meaning or function
- Any bus the lines can be classified into three functional groups: **data**, **address**, and **control** lines
- There may be **power distribution lines** that supply power to the attached modules

# DATA BUS

- The data lines provide a path for moving data among system modules
  - These lines, collectively, are called the *data bus*
- The data bus may consist of 32, 64, 128, or even more separate lines, the number of lines being referred to as the *width* of the data bus
- Because each line can carry only 1 bit at a time, the number of lines determines how many bits can be transferred at a time
- **Width:** If the data bus is 32 bits wide and each instruction is 64 bits long, then the processor must access the memory module twice during each instruction cycle

# ADDRESS BUS

- The address lines are used to designate the source or destination of the data on the data bus
- For **example**, if the processor wishes to read a word (8, 16, or 32 bits) of data from memory, it puts the address of the desired word on the address lines
- The width of the address bus determines the maximum possible memory capacity of the system.
- Typically, the higher-order bits are used to select a particular module on the bus, and the lower-order bits select a memory location or I/O port within the module
- For **example**, on an 8-bit address bus, address 01111111 and below might reference locations in a memory module (module 0) with 128 words of memory, and address 10000000 and above refer to devices attached to an I/O module (module 1)

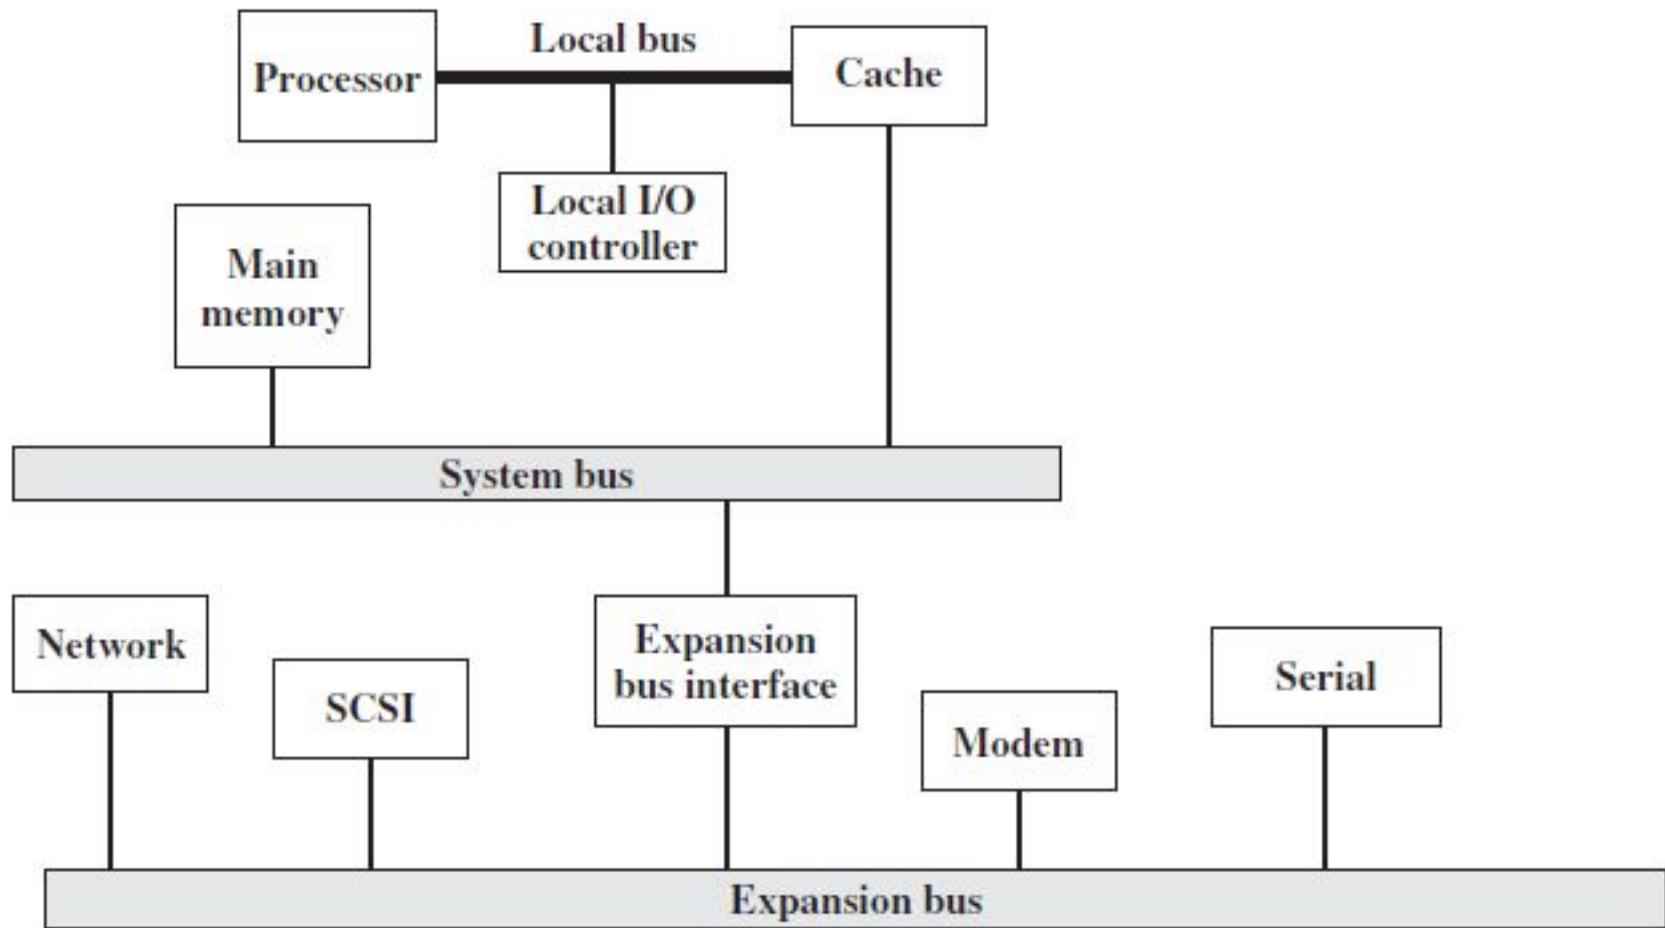
# CONTROL BUS

- The control lines are used to control the access to
- Control signals transmit both command and timing information among system modules
  - Timing signals indicate the validity of data and address information
  - Command signals specify operations to be performed
- Typical control lines include:
  - Memory write, Memory read, I/O write, I/O read, Transfer ACK, Bus request, Bus grant, Interrupt request, Interrupt ACK, Clock, Reset

# Multiple-Bus Hierarchies

- Motivated by two major concerns:
  - 1) In general, the more devices attached to the bus, the greater the bus length and hence the greater the propagation delay
  - 2) The bus may become a bottleneck as the aggregate data transfer demand approaches the capacity of the bus
- Most computer systems use multiple buses, generally laid out in a hierarchy

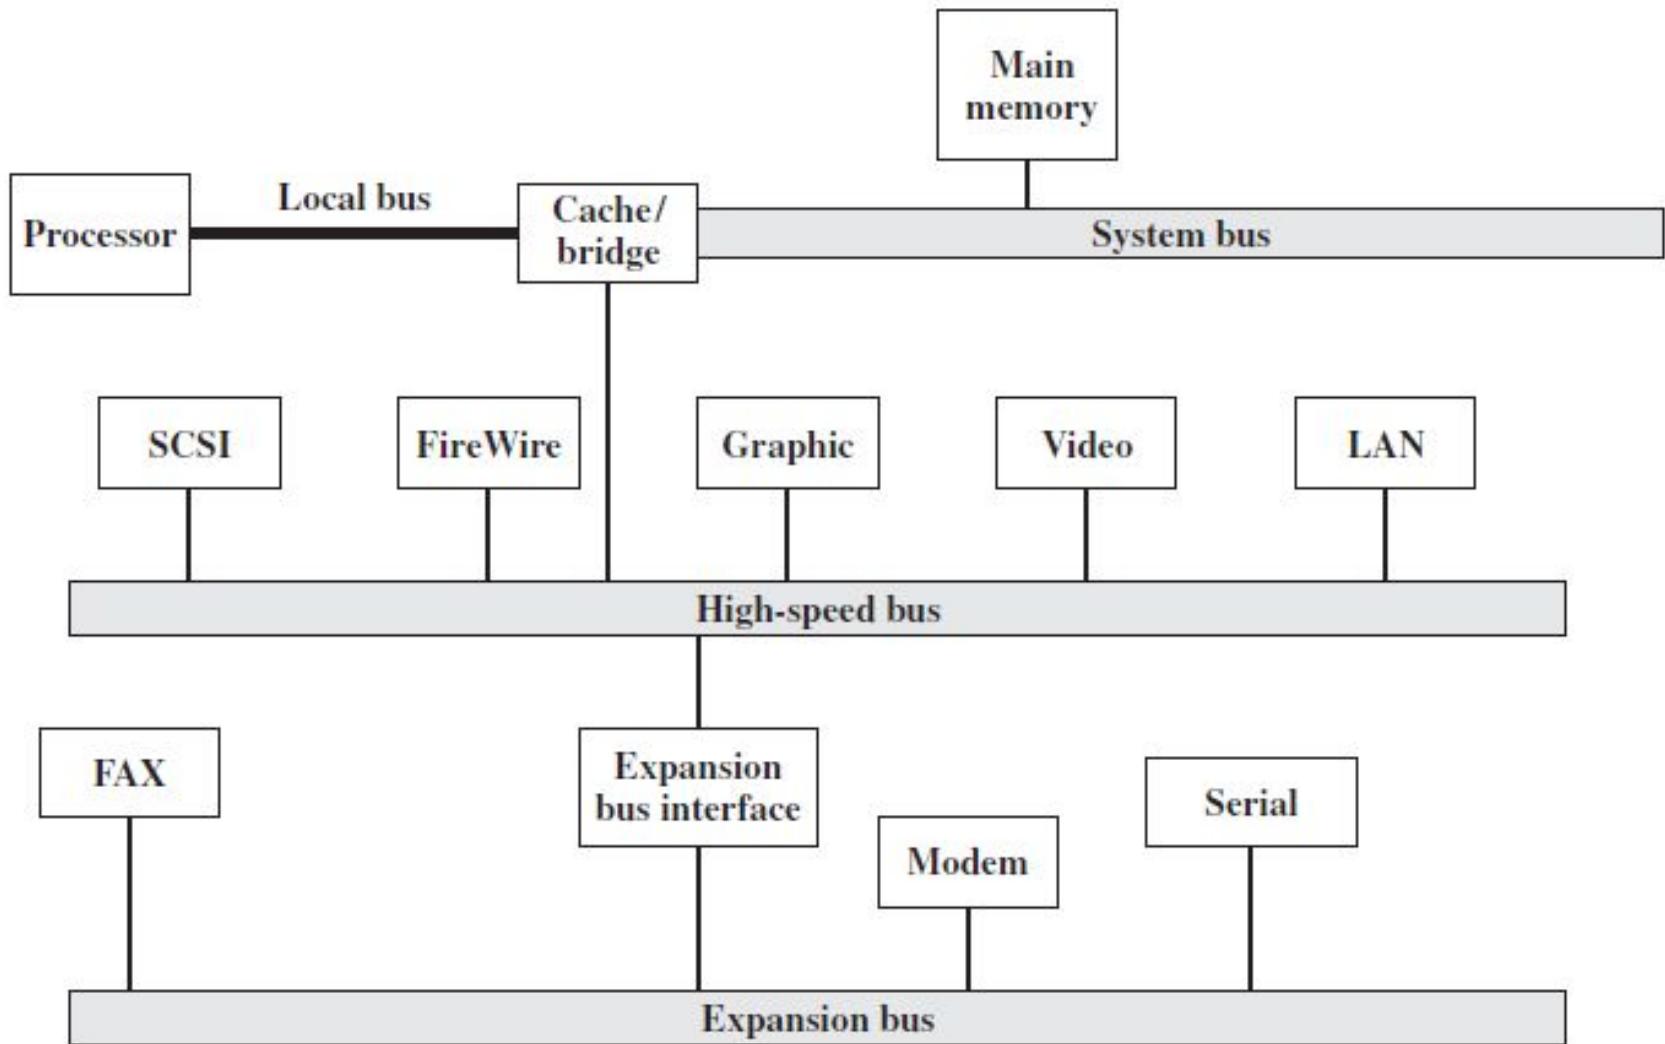
# Traditional BUS Architecture



# Traditional BUS Architecture

- Network connections include local area networks (LANs) such as a 10-Mbps Ethernet and connections to wide area networks (WANs) such as a packet-switching network
- SCSI (small computer system interface) is itself a type of bus used to support local disk drives and other peripherals
- A serial port could be used to support a printer or scanner

# High-Performance



# High-Performance

- **Architecture**: There is a local bus that connects the processor to a cache controller, which is in turn connected to a system bus that supports main memory
- The cache controller is integrated into a bridge, or buffering device, that connects to the high-speed bus
- This bus supports connections to high-speed LANs, such as Fast Ethernet at 100 Mbps, video and graphics workstation controllers, as well as interface controllers to local peripheral buses, including SCSI and FireWire
- The latter is a high-speed bus arrangement specifically designed to support high-capacity I/O devices
- Lower-speed devices are still supported off an expansion bus, with an interface buffering traffic between the expansion bus and the high-speed bus

# High-Performance Architecture

- The advantage of this arrangement is that the high-speed bus brings high demand devices into closer integration with the processor and at the same time is independent of the processor
- Thus, differences in processor and high-speed bus speeds and signal line definitions are tolerated
- Changes in processor architecture do not affect the high-speed bus, and vice versa

# Elements of Bus Design

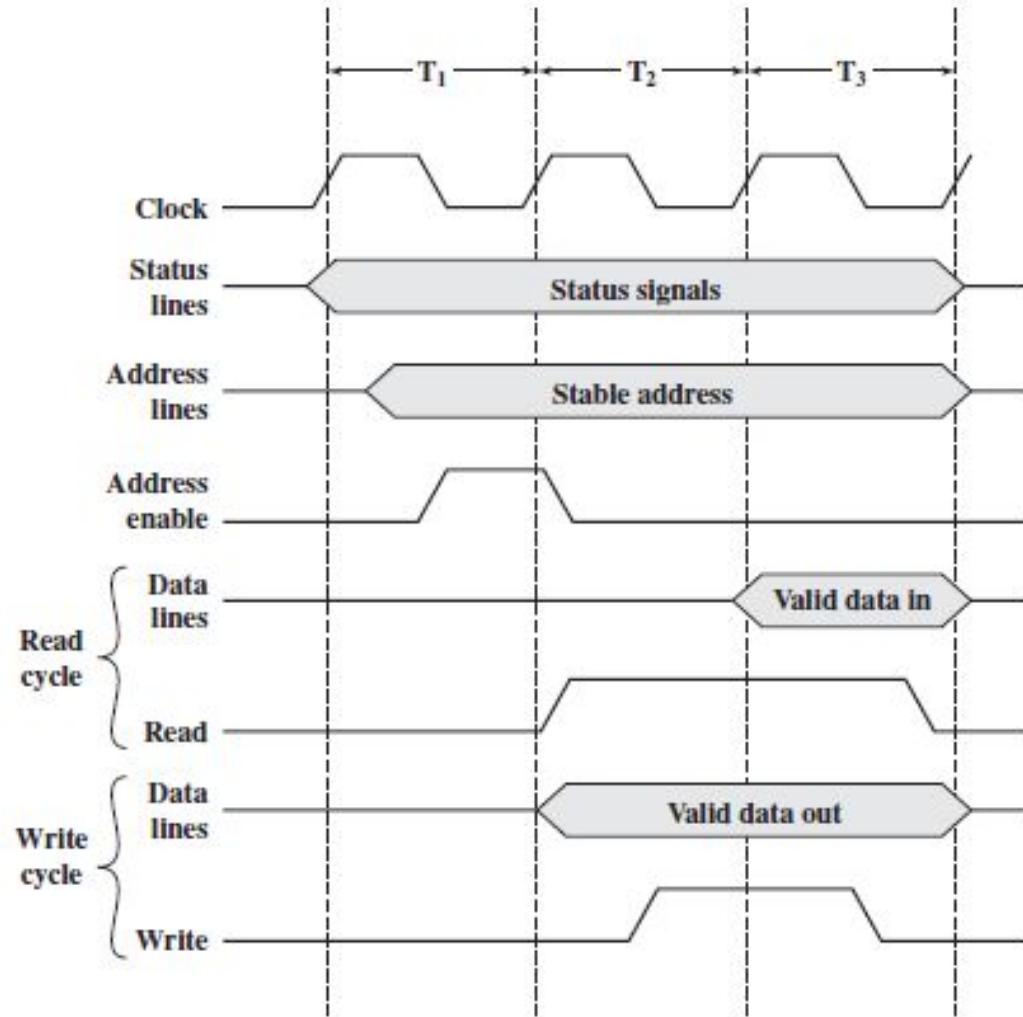
Type	Bus Width
Dedicated	Address
Multiplexed	Data
Method of Arbitration	Data Transfer Type
Centralized	Read
Distributed	Write
Timing	
Synchronous	Read-modify-write
Asynchronous	Read-after-write
	Block

# Method of Arbitration

- Roughly **Two** categories: Centralized and Distributed
- In a **centralized** scheme, a single hardware device, referred to as a *bus controller* or *arbiter*, is responsible for allocating time on the bus
- The device may be a separate module or part of the processor
- In a **distributed** scheme, there is no central controller
- Rather, each module contains access control logic and the modules act together to share the bus
- With both methods of arbitration, the purpose is to designate one device, either the processor or an I/O module as master

# Timing

- With **synchronous timing**, the occurrence of events on the bus is determined by a clock
- The bus includes a clock line upon which clock transmits a regular sequence of alternating 1s and 0s of equal duration



# Timing

- With **asynchronous timing**, the occurrence of one event on a bus follows and depends on the occurrence of a previous event

